Recap

- Sqoop

- HBase

Agenda for today

Hadoop

| Processing nature	Tool
Batch processing on structured data	Hive
Ad-hoc analysis on structured data	Impala
Machine Learning	Apache Mahout
Graph processing	Apache Giraph
Stream processing	Apache Storm

Spark

Processing nature	Spark
Batch processing on structured data	SparkSQL
Ad-hoc analysis on structured data	SparkSQL
Machine Learning	MlLib
Graph processing	Graphax
Stream processing	Spark Streaming

Why Spark?

- ▶ Faster

Logistic regression in Hadoop and Spark

Why Spark: cont...

- ▶ Ease of use
- ❑ Support for multiple languages
- ❑ REPL for development and ad-hoc analysis
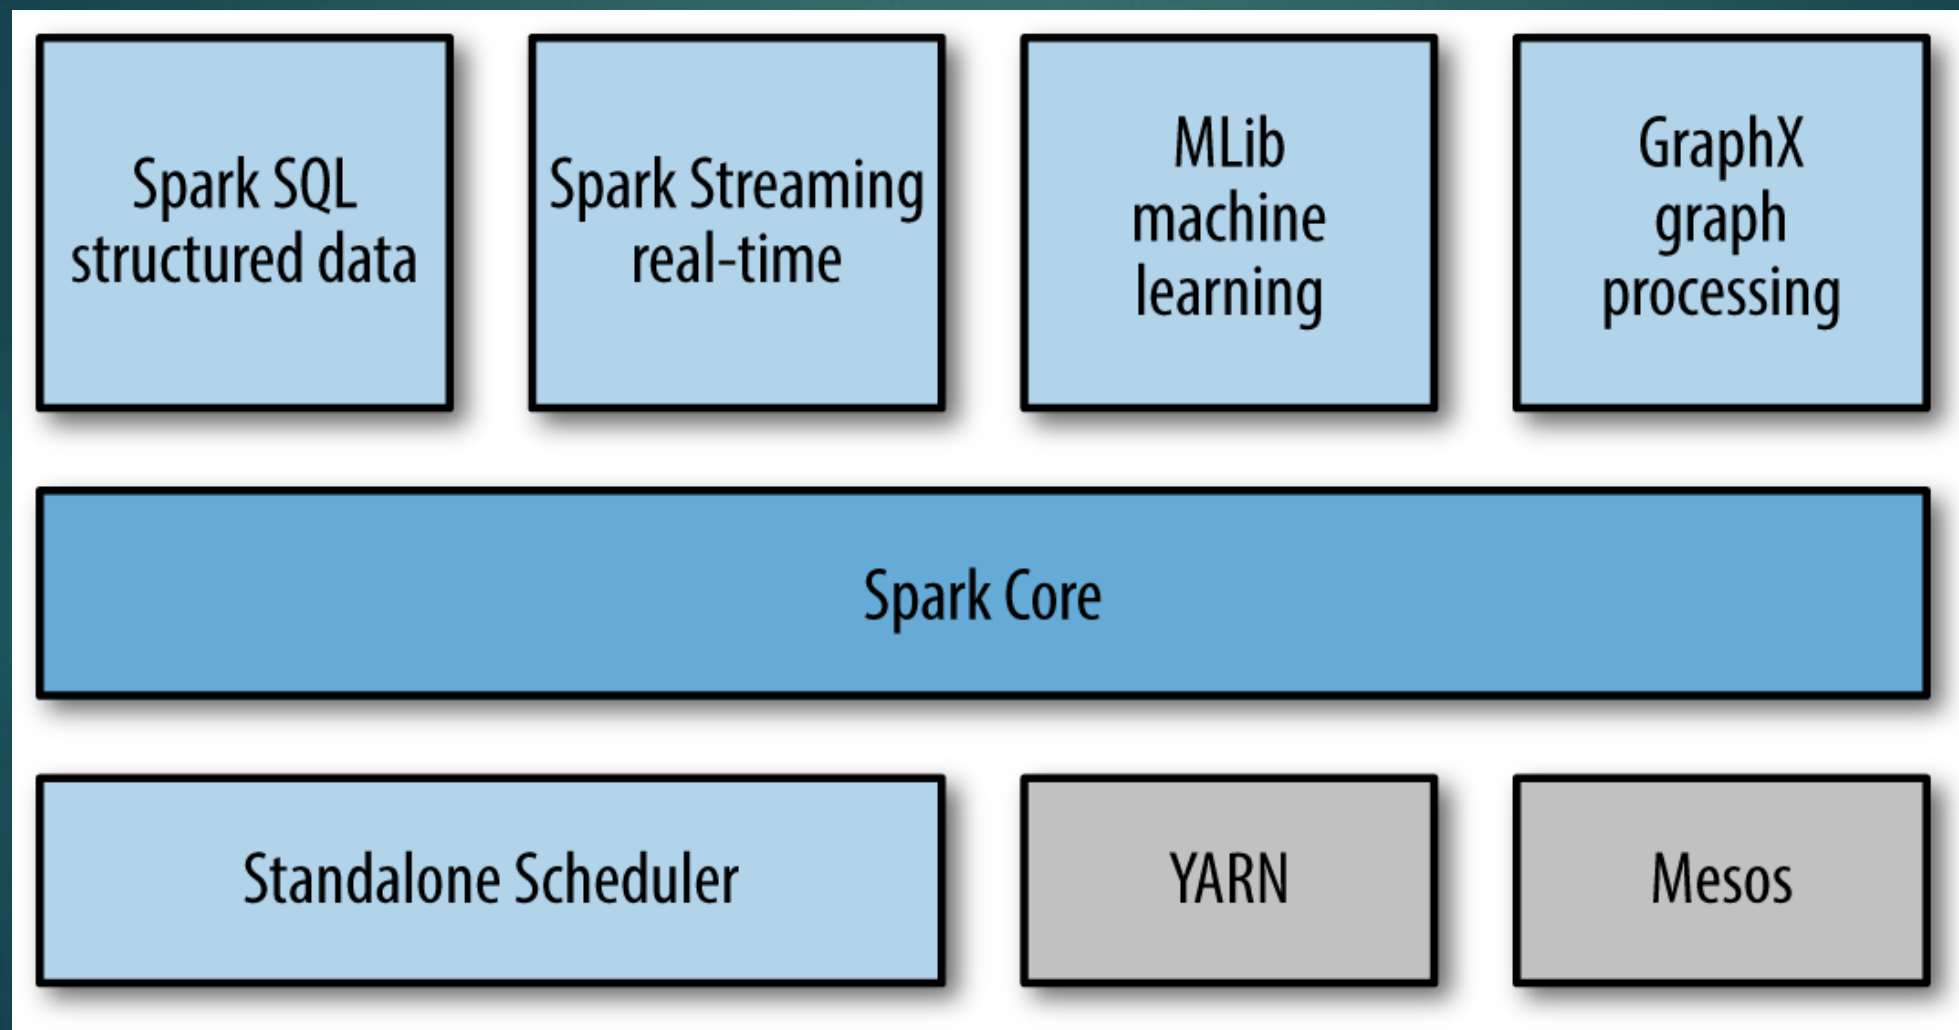- ❑ Fewer lines of code

Why Spark: cont…

- ▶ Inter operability with other platforms
- ❑ Hadoop
- ❑ Mesos
- ❑ Hbase
- ❑ Cassandra

Components

RDD

- **R**esilient **D**istributed **D**ataset

  Fault tolerant

  Distributed across multiple processes

  Source could be a file or program generated

- Immutable collection of elements, partitioned across multiple processes to operate in parallel

RDD Operation: Transformation

- Evaluated lazily

- Can be applied on any RDD

- Generates another RDD as result

- Example: map, flatMap, filter, reduceByKey…

RDD Operation: Action

- ▶ Call for evaluation of complete DAG

- ▶ Can be applied on any RDD

- ▶ Generates result on driver program
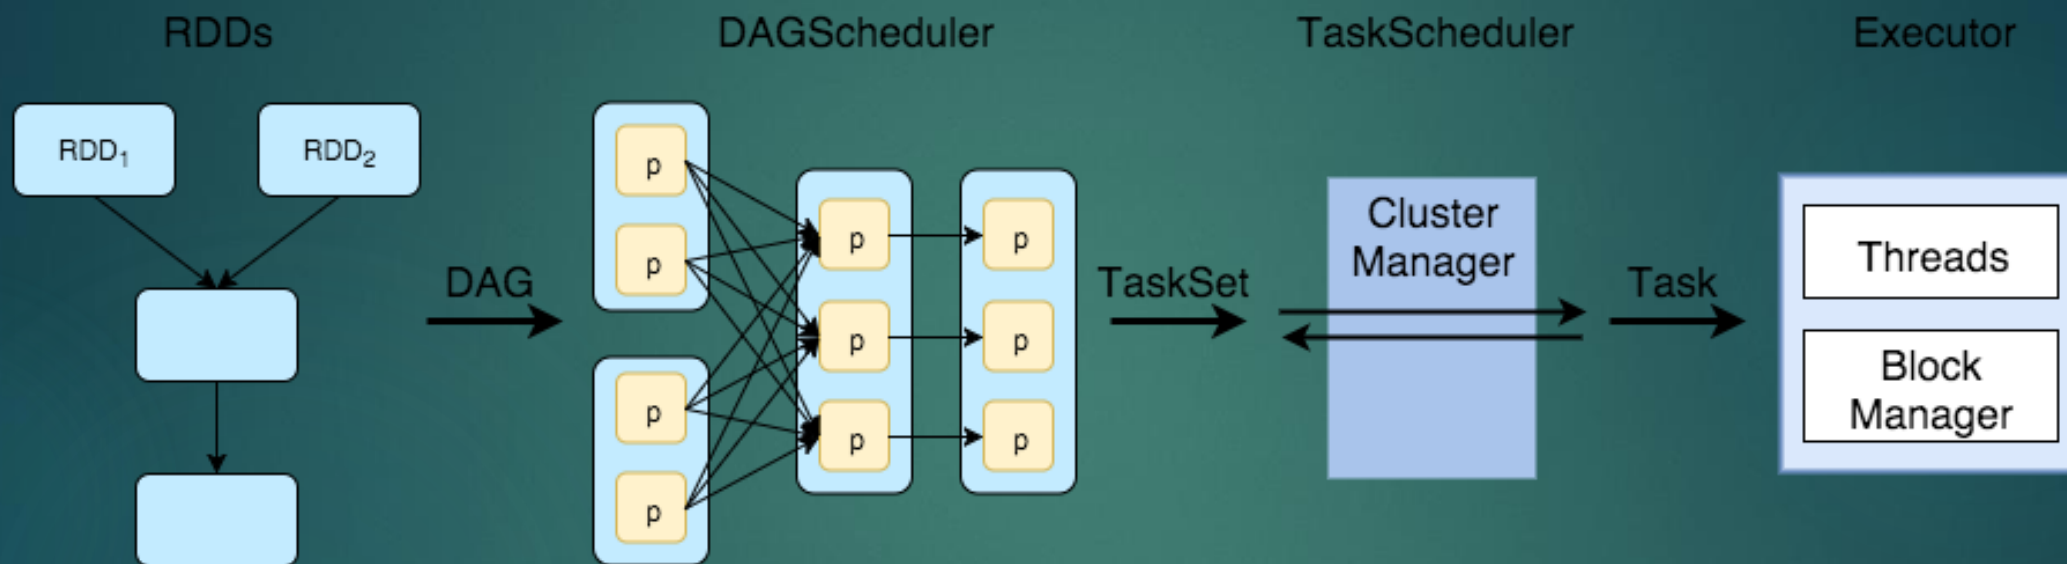
- ▶ Example: count, take, saveAsTextFile, collect…

DAG

- Directed Acyclic Graph prepared on RDD

- DAG scheduler prepares stages and tasks

- Tasks within a stage will be executed when stage is ready to execute

- Shuffle operation is the stage boundary

Execution and Coordination

Launch spark shell

- Standalone:

./bin/spark-shell --master spark://IP:PORT

- YARN

./bin/spark-shell --master yarn

- Mesos

./bin/spark-shell --master mesos://host:5050

Submit Application

▶ Scala/Java

spark-submit --class SparkWordCount --master local --deploy-mode client --executor-memory 1g --name wordcount --conf "spark.app.id=wordcount" sparkwordcount.jar <other parameters to JAR file>

▶ Python

spark-submit --master yarn --deploy-mode client --executor-memory 1g --name wordcount --conf "spark.app.id=wordcount" wordcount.py <Other parameters>

Language comparison matrix

Metrics	Scala	Java	Python	R
Type	Compiled	Compiled	Interpreted	Interpreted
JVM based	Yes	Yes	No	No
Verbosity	Less	More	Less	Less
Code Length	Less	More	Less	Less
Productivity	High	Less	High	High
Scalability	High	High	Less	Less
OOPS Support	Yes	Yes	Yes	Yes

REPL

- Scala
spark-shell

- Python
pyspark

Why SparkSQL?

- ▶ Integrated

context = HiveContext(sc)
results = context.sql(
  *"SELECT * FROM people")*
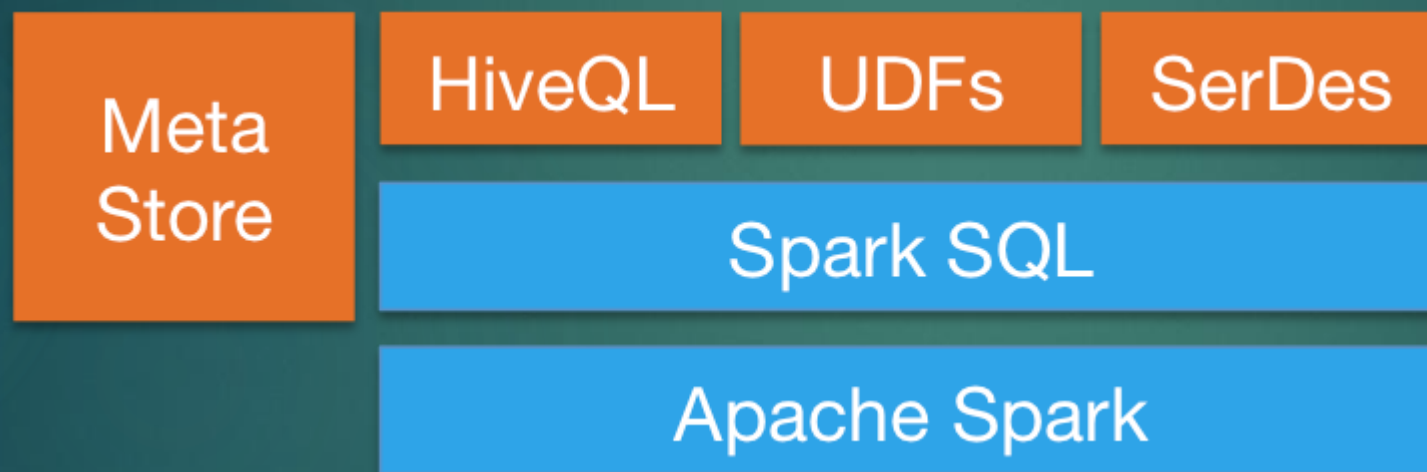*names = results.map(lambda p: p.name)*

Why SparkSQL: cont…

▶ Uniform Data access

```
context.jsonFile("filename.json")
  .registerTempTable("json")
results = context.sql(
  """SELECT *
    FROM people
    JOIN json …""")
```

Why SparkSQL: cont…

▶ Hive Integration

Why SparkSQL: cont…

- ▶ Standard Connectivity

BI Tools

...

JDBC / ODBC

Spark SQL

Introduction

▶ Relationship with Shark project

    ❑ Shark had limited integration with Spark

    ❑ Hive optimizer was not best fit for Spark

    ❑ Spark SQL reused Hive data loading as well as in-memory column storage features of Shark

    ❑ Additionally, introduced RDD-aware optimizer and rich language interface

Data holders

- RDD

- Dataframe: RDD with schema

- Dataset:
  introduced in 1.6 version
  provides strong type over RDD

Data source

- ▶ Hive existing table

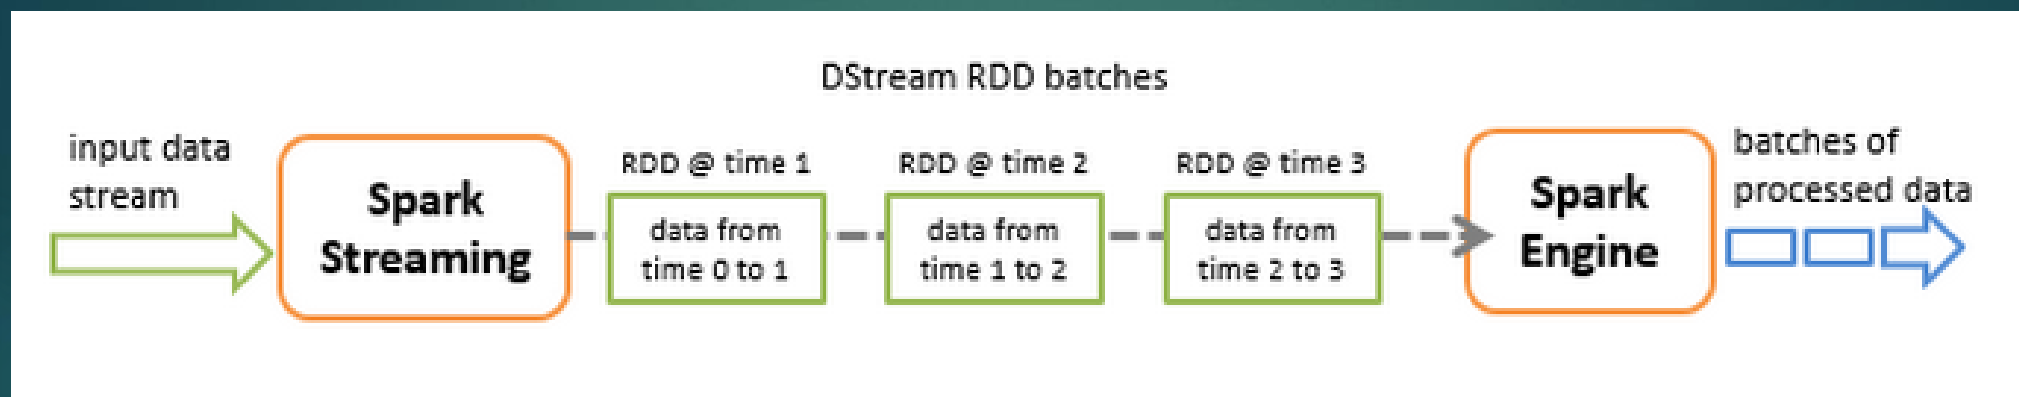- ▶ Structured files. Json file for example

- ▶ RDD

Hive integration

- ▶ Can use most of the SQL features available in Hive

- ▶ Insert data through Spark and read in Hive

- ▶ Executes DDL statements

- ▶ Refers Hive metastore for metadata

DStream

DStream RDD batches

input data stream → **Spark Streaming** — RDD @ time 1: data from time 0 to 1 — RDD @ time 2: data from time 1 to 2 — RDD @ time 3: data from time 2 to 3 → **Spark Engine** → batches of processed data

Source

- Kafka
- Flume
- HDFS/S3
- Kinesis
- Twitter

Storage/Target

- ▶ HDFS

- ▶ Databases

- ▶ Dashboard

Program flow

- ▶ Set streaming context

- ▶ Define source for the streaming context

- ▶ Apply all transformations of Dstream

- ▶ Start the streaming context

Examples

- https://github.com/apache/spark/tree/master/examples/src/main/scala/org/apache/spark/examples

Persist/Cache data

- Helpful to reuse the same dataset

- Multiple storage levels:
https://spark.apache.org/docs/latest/rdd-programming-guide.html

- How to check current storage level:
  <Object name>.getStorageLevel

Further studies

- https://www.cloudera.com/documentation/enterprise/5-6-x/PDF/cloudera-spark.pdf

- https://databricks.com/product/getting-started-guide/quick-start

- Cloud hosted community spark setup
  https://community.cloud.databricks.com/

References

- http://spark.apache.org/docs/1.3.0/cluster-overview.html
- Hadoop: the definitive guide 4th edition