



Recap

2

- ▶ Sqoop

- ▶ HBase

Agenda for today

3



Hadoop

Processing nature	Tool
Batch processing on structured data	Hive
Ad-hoc analysis on structured data	Impala
Machine Learning	Apache Mahout
Graph processing	Apache Giraph
Stream processing	Apache Storm

Spark


5

Processing nature	Spark
Batch processing on structured data	SparkSQL
Ad-hoc analysis on structured data	SparkSQL
Machine Learning	MLlib
Graph processing	Graphax
Stream processing	Spark Streaming

Why Spark?

6

► Faster



Logistic regression in Hadoop and Spark

Why Spark: cont...

7

- ▶ Ease of use
 - ❑ Support for multiple languages
 - ❑ REPL for development and ad-hoc analysis
 - ❑ Fewer lines of code

Why Spark: cont...

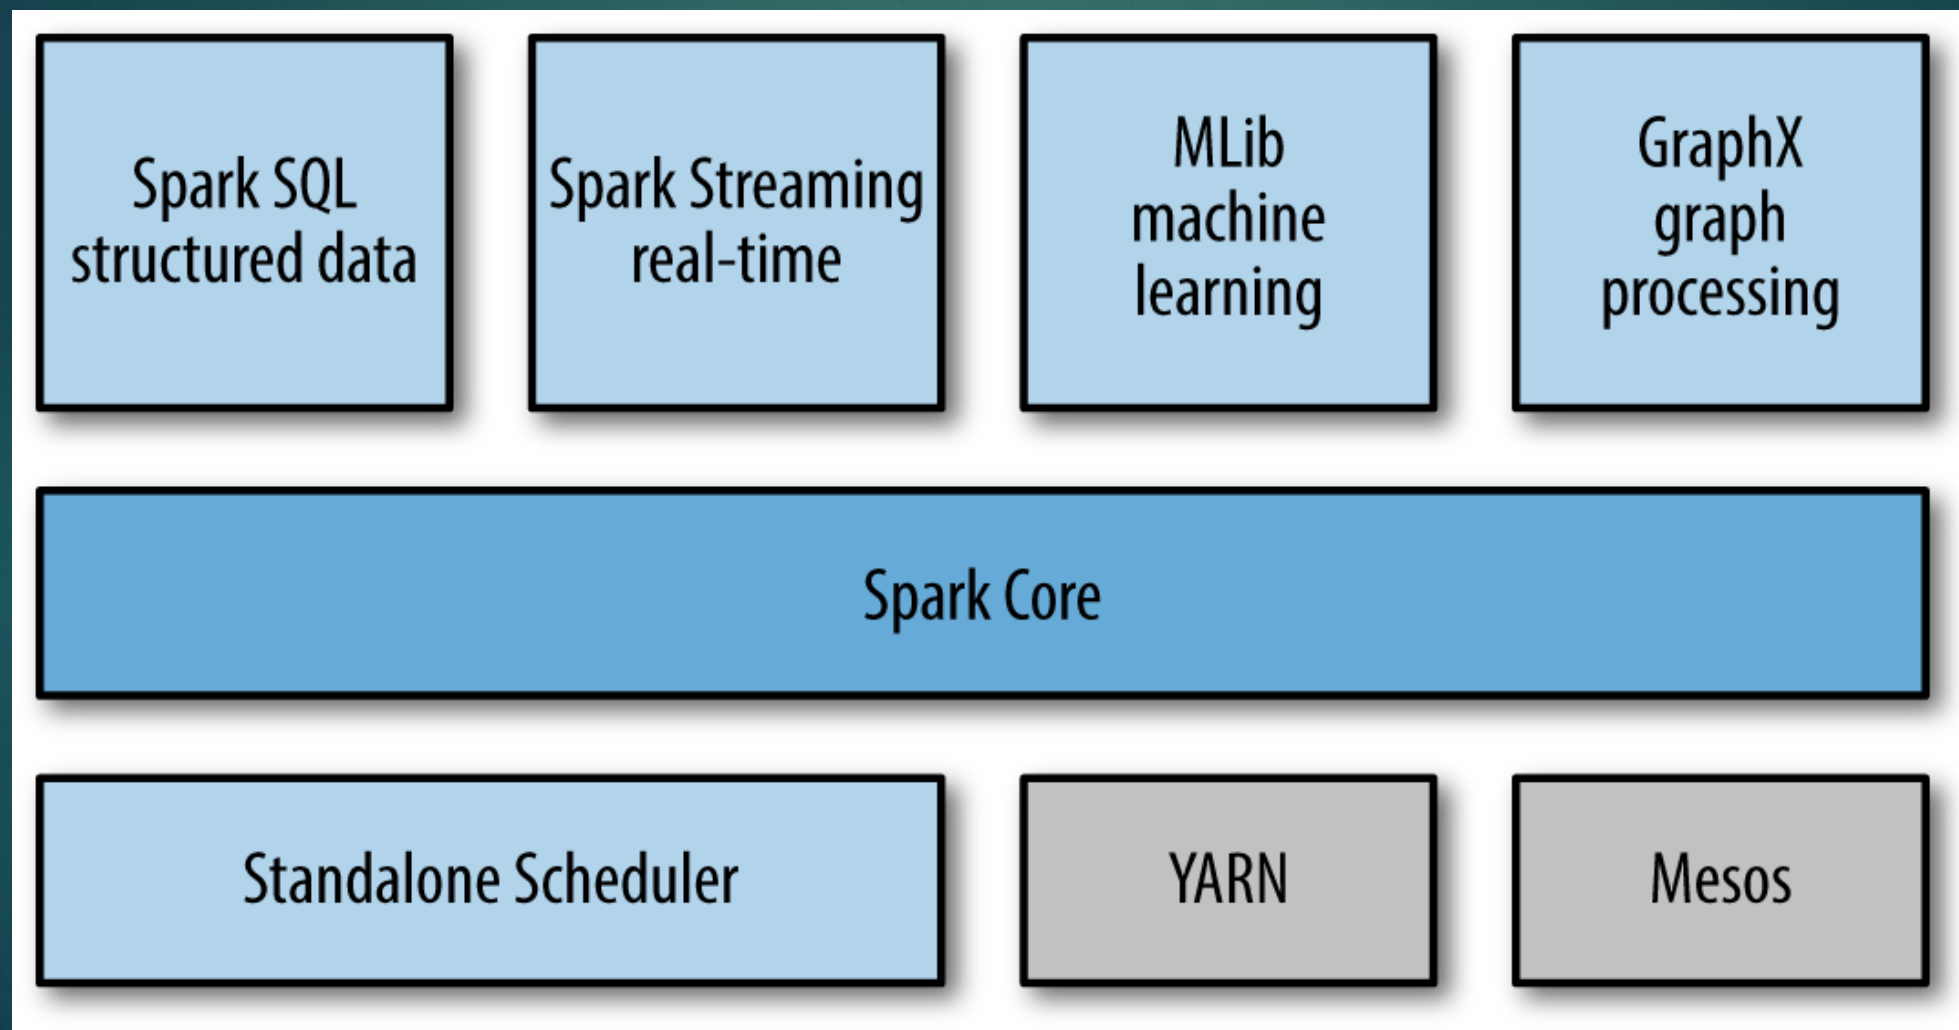
8

► Inter operability with other platforms

- ❑ Hadoop
- ❑ Mesos
- ❑ Hbase
- ❑ Cassandra

Components

9



- ▶ **R**esilient **D**istributed **D**ataset

 - Fault tolerant

 - Distributed across multiple processes

 - Source could be a file or program generated

- ▶ Immutable collection of elements, partitioned across multiple processes to operate in parallel

RDD Operation: Transformation

11

- ▶ Evaluated lazily
- ▶ Can be applied on any RDD
- ▶ Generates another RDD as result
- ▶ Example: map, flatMap, filter, reduceByKey...

RDD Operation: Action

12

- ▶ Call for evaluation of complete DAG
- ▶ Can be applied on any RDD
- ▶ Generates result on driver program
- ▶ Example: count, take, saveAsTextFile, collect...

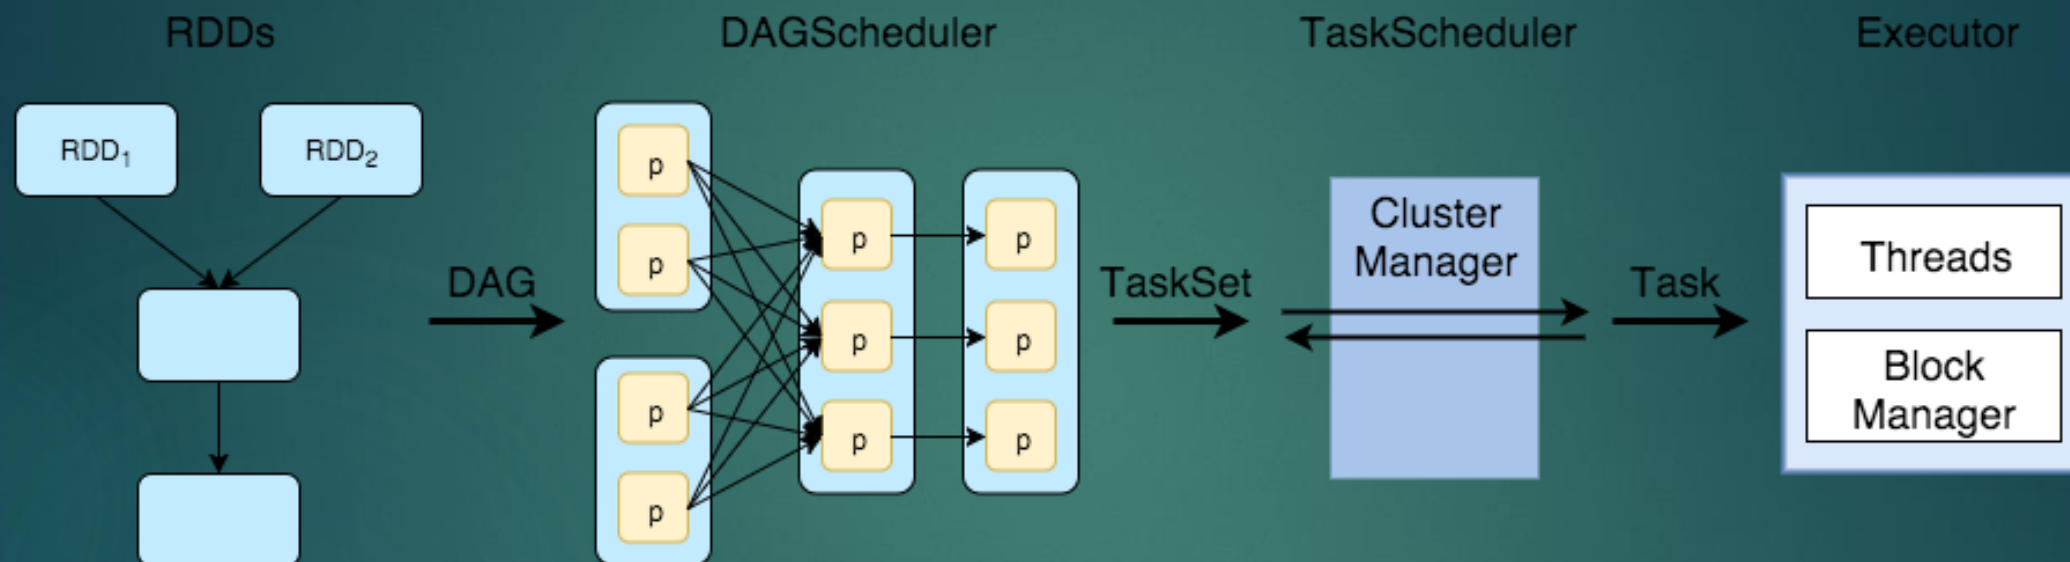
DAG

13

- ▶ Directed Acyclic Graph prepared on RDD
- ▶ DAG scheduler prepares stages and tasks
- ▶ Tasks within a stage will be executed when stage is ready to execute
- ▶ Shuffle operation is the stage boundary

Execution and Coordination

14



Launch spark shell

15

- ▶ Standalone:

```
./bin/spark-shell --master spark://IP:PORT
```

- ▶ YARN

```
./bin/spark-shell --master yarn
```

- ▶ Mesos

```
./bin/spark-shell --master mesos://host:5050
```

Submit Application

16

► Scala/Java

```
spark-submit --class SparkWordCount --master local --  
deploy-mode client --executor-memory 1g --name  
wordcount --conf "spark.app.id=wordcount"  
sparkwordcount.jar <other parameters to JAR file>
```

► Python

```
spark-submit --master yarn --deploy-mode client --  
executor-memory 1g --name wordcount --conf  
"spark.app.id=wordcount" wordcount.py <Other  
parameters>
```

Language comparison matrix

17

Metrics	Scala	Java	Python	R
Type	Compiled	Compiled	Interpreted	Interpreted
JVM based	Yes	Yes	No	No
Verbosity	Less	More	Less	Less
Code Length	Less	More	Less	Less
Productivity	High	Less	High	High
Scalability	High	High	Less	Less
OOPS Support	Yes	Yes	Yes	Yes

REPL

18

► Scala
spark-shell

► Python
pyspark

The logo for Spark SQL. It features the word "Spark" in a bold, black, sans-serif font. An orange, five-pointed star is positioned above the letter "k". To the right of "Spark" is the text "SQL" in a black, sans-serif font, all in uppercase.

Spark SQL

Spark SQL: Why?

20

► Integrated

```
context = HiveContext(sc)
results = context.sql(
    "SELECT * FROM people")
names = results.map(lambda p: p.name)
```


Spark SQL: Why?

21


► Uniform Data access

```
context.jsonFile("filename.json")  
  .registerTempTable("json")  
results = context.sql(  
  """SELECT *  
    FROM people  
    JOIN json ...""")
```

Spark SQL: Why?

22

► Hive Integration



Spark SQL: Why?

23

► Standard Connectivity

BI Tools

...

JDBC / ODBC

Spark SQL

Introduction

24

- ▶ Relationship with Shark project
 - ❑ Shark had limited integration with Spark
 - ❑ Hive optimizer was not best fit for Spark
 - ❑ Spark SQL Reused Hive data loading as well as in-memory column storage features of Shark
 - ❑ Additionally, introduced RDD-aware optimizer and rich language interface

Data holders

25

- ▶ RDD
- ▶ Dataframe: RDD with schema
- ▶ Dataset:
introduced in 1.6 version
provides strong type over RDD

Data source

26

- ▶ Hive existing table
- ▶ Structured files. Json file for example
- ▶ RDD

Hive integration

27

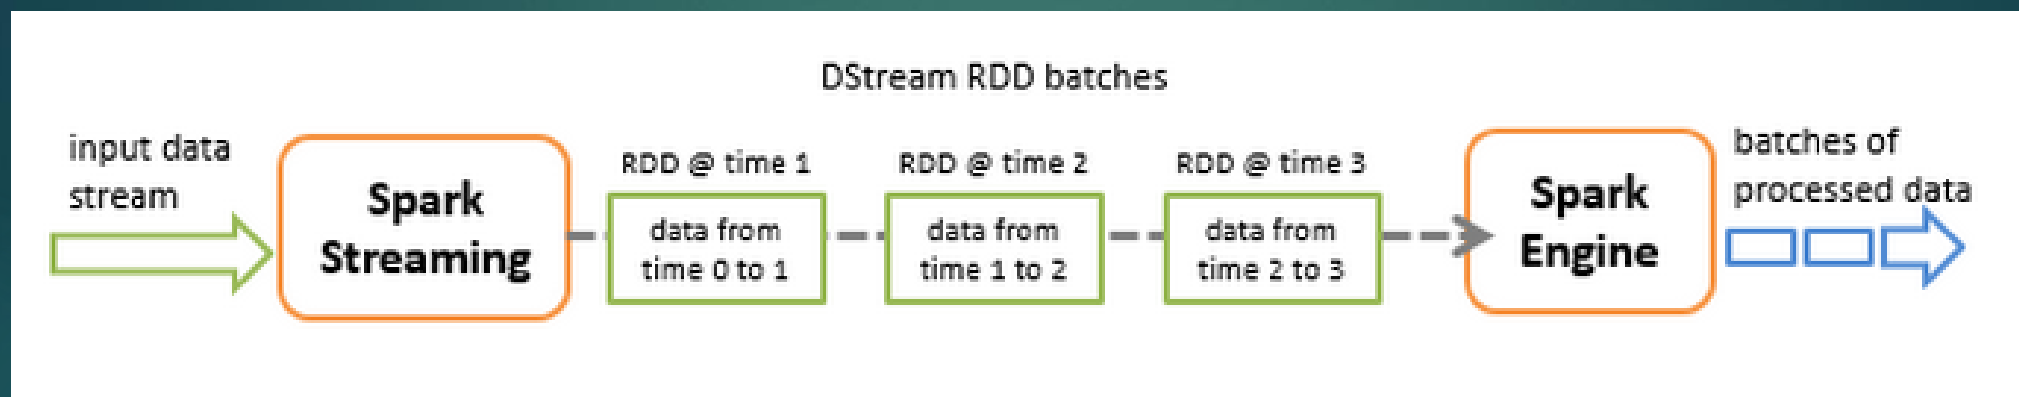
- ▶ Can use most of the SQL features available in Hive
- ▶ Insert data through Spark and read in Hive
- ▶ Executes DDL statements
- ▶ Refers Hive metastore for metadata



Spark Streaming

DStream

29



Source

30

- ▶ Kafka
- ▶ Flume
- ▶ HDFS/S3
- ▶ Kinesis
- ▶ Twitter

Storage/Target

31

- ▶ HDFS
- ▶ Databases
- ▶ Dashboard

Program flow

32

- ▶ Set streaming context
- ▶ Define source for the streaming context
- ▶ Apply all transformations of Dstream
- ▶ Start the streaming context

Examples

33

- ▶ <https://github.com/apache/spark/tree/master/examples/src/main/scala/org/apache/spark/examples>

Persist/Cache data

34

- ▶ Helpful to reuse the same dataset

- ▶ Multiple storage levels:

<https://spark.apache.org/docs/latest/rdd-programming-guide.html>

- ▶ How to check current storage level:

`<Object name>.getStorageLevel`

Further studies

35

- ▶ <https://www.cloudera.com/documentation/enterprise/5-6-x/PDF/cloudera-spark.pdf>
- ▶ <https://databricks.com/product/getting-started-guide/quick-start>
- ▶ Cloud hosted community spark setup
<https://community.cloud.databricks.com/>

References

36

- ▶ <http://spark.apache.org/docs/1.3.0/cluster-overview.html>
- ▶ Hadoop: the definitive guide 4th edition